

On the Ground Validation of Online Diagnosis with Twitter and Medical Records

Todd Bodnar^{*}

Pennsylvania State University
Center for Infectious Disease
Dynamics and Department
of Biology
meme@psu.edu

Conrad S Tucker

Pennsylvania State University
Department of Engineering
Design and Industrial and
Manufacturing Engineering
conrad.tucker@psu.edu

Victoria C Barclay

Pennsylvania State University
Center for Infectious Disease
Dynamics and Department
of Biology
vickistar50@gmail.com

Marcel Salathé

Pennsylvania State University
Center for Infectious Disease
Dynamics and Department
of Biology
salathe@psu.edu

ABSTRACT

Social media has been considered as a data source for tracking disease. However, most of the analysis has been on the population level. Here we develop a novel system of influenza diagnosis based on publicly available Twitter data. We train our system on cases where the Twitter user has been officially diagnosed by medical professionals. We find that about half ($17/35 = 48.57\%$) of the users in our sample that were sick explicitly discuss their disease on Twitter. By developing a meta classifier based off of text analysis, anomaly detection, and social network analysis, we are able to diagnose an individual with greater than 99% accuracy even if she does not discuss her health.

Categories and Subject Descriptors

I.2.1 [Artificial Intelligence]: Applications and Expert Systems—*Medicine and Science*

General Terms

Experimentation, Validation

Keywords

Twitter, Validation, Digital Epidemiology, Remote Diagnosis

1. INTRODUCTION

Disease surveillance systems – traditionally relying on reports from medical practitioners – are an important part of

disease control. However, these traditional surveillance systems are often costly and slow to respond[8, 3, 13]. The widespread adoption of the Internet by the general public has provided opportunities for the development of novel disease surveillance methods. Compared to traditional systems, where data is provided by medical diagnosis, these new systems provide either semi-automatic – through long term self reporting systems[10, 16] – or fully automatic – through data mining search queries or social media[1, 4, 5, 2, 11] – disease surveillance. While these methods are cheaper, faster and cover a larger number of individuals than traditional systems, one can be less confident about their results than the results from a system based on professional diagnosis. In this paper, we develop a system that performs long term surveillance on Twitter users with methods trained on professionally diagnosed data that combines the advantages of all three of these systems.

Previous work with data mining social media has focused on methods to replicate the patterns found in traditional surveillance networks[1, 4, 5]. However, these methods have several limitations. First, they generally do not differentiate between an individual with an illness and an individual that is worried about an illness; which may have resulted in a predicted influenza rate that was much higher than the actual 2013 influenza rate [1, 2, 11, 9]. Second, these methods cannot be extended to areas without a previous surveillance network. Finally, these methods are fundamentally incapable of detecting diseases that do not show strong spatial-temporal patterns such as mental illness, obesity or Parkinson's disease. Instead of top-down methods to measure levels of disease in a population, we approach this problem from the bottom-up. This addresses all three of these issues: we only diagnose individuals that are likely to have the disease, and not just interested in the disease; we do not require previous data when applying these methods to new problems or locations; and these methods can easily generalize to diseases that do not show strong patterns because we focus on an individual level.

^{*}Corresponding author

Systems, such as InfluenzaNet or Flu Near You, use self-reported symptoms to diagnose an individual also work from a bottom-up approach.[10, 16] These systems have the potential to be better than traditional surveillance systems because they update in near-real-time and can detect cases even when the user has not gone to their doctor. These systems require the user to sign up which allows for long term studies which are not normally able to be done with Tweets or search queries. However this reduces the number of users studied compared to data mining approaches. For example, Flu Near You had a total of 9,456 users report during the week ending 29 December 2013. Marquet et al. [10] have shown a large drop out rate with only 53% users participating for five or more weeks. While this amount of data is sufficient for many purposes, a system which is based on Twitter’s millions of active users would open the door to more questions.

We develop such a system as follows. In section 2 we describe the collection of professional diagnosis of Twitter users and collect their Twitter information. In section 3 we consider extracting textual information from Tweets as a method for diagnosing influenza. Previous work has focused on this area. Additionally, we consider other methods for detection. In section 4 we consider anomalies in a user’s Tweeting behavior as a signal for diagnosing influenza. In section 5 we extend these methods to other users on a persons social network to diagnose the original person. In section 6 we aggregate the results of the previous classifiers to develop a more accurate meta-classifier.

2. DATA COLLECTION

2.1 Medical Records

We received information from the Pennsylvania State University’s Health Services about 104 individuals that were diagnosed with influenza by a medical professional during the 2012-2013 Influenza season (see figure 1). Additionally we collect information from 122 individuals that were *not* diagnosed as a control set. The participants were mostly students (72% were between the ages of 18 and 22) and slightly more female than expected (133/226 \approx 58.8%.) Data collection was approved through the Pennsylvania State University’s IRB (#41345.) Additionally we recieved the Twitter handles from 119 of these individuals.

2.2 Twitter Records

While we received a total of 119 Twitter accounts, 15 were discarded because the associated accounts were either non-existent, banned, or private. For each of the remaining 104 accounts, we pulled their profile information, their friends and followers information, their most recent 3000 tweets, and their friends profiles and tweets. Some users did not tweet during the month that they were sick; we kept those accounts as part of the control group. We were limited to the most recent 3000 tweets by Twitter’s time line query, but this only effected two accounts – both of which posted multiple times per hour and were thrown out because we could only look back a few days.

We collected data by calling the Twitter API on the user account that we queried the longest time ago. Tweets, profile and follower information queries have separate rate limits




Figure 1: The rate of professionally diagnosed Influenza cases during the 2012-2013 season in our sample.

and were collected in parallel. The 104 seed accounts collected above were given higher priority over their friends and followers. In total, we collected 37,599 tweets from the seed accounts and 30,950,958 tweets from 913,082 accounts that they either followed or were followed by.

3. TEXT BASED SIGNALS

In this section, we consider diagnosis based on the content of a user’s tweets. Such analysis can be approached by keyword analysis, where the presence of absence of a keyword predicts disease, or through text classification, where the tweets are classified as being about disease or not about disease. We begin by dividing the tweets into two sets: tweets that were posted the same month that a user was sick and tweets that were posted other times. We find a total of 1609 tweets from 35 users in the first category.

First we go the route of defining a set of keywords that are positive signals of influenza. We chose {flu, influenza, sick, cough, cold, medicine, fever} as our set of keywords. These keywords include the names and symptoms of the illness in addition to “medicine” and serve as a set of keywords that may have been chosen by a domain expert. Of these seven keywords, we find a significant effect in 6 of the keywords during months when the user had influenza (See table 1). Additionally, we try algorithmically selecting keywords by first finding the 12,393 most common keywords in the data set. We then rank them based off of information gain and choose the top 10, 100 or 1000 keywords from the list. In both of these cases, we pre-process the data by tokenizing the text on spaces, tabs and line breaks and the characters “.,;:’()?!/\” and remove stop words¹, perform Porter stemming [12, 17] and convert the text to lower case. We then use the occurrence or absence of these keywords as features for classification. We use Naive Bayes, random forest, J48 (a Java implementation of C4.5), logistic regression and support vector machines to classify a user as being sick in a given month or not (see figure 2).

Additionally, we hand rate all 1609 tweets that were posted by individuals during the time of their illness for informa-

¹Stop words taken from Weka’s stop list version 3.7.10.

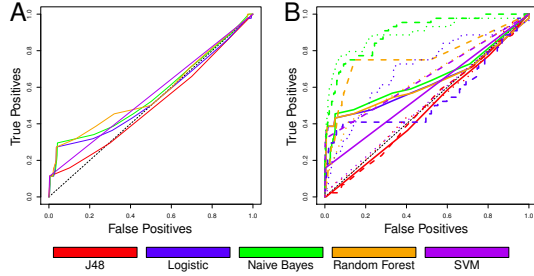


Figure 2: The ROC of classifiers that use hand chosen keywords (a) and algorithmically chosen keywords (b) to determine if an individual is ill. The top 10 (solid line), 100 (dashed line) and 1000 (dotted line) were selected as the features.

Word	Total	Odds Ratio	Significance
flu	25	40.14	<0.0001
influenza	1	0.00	0.8325
sick	128	5.22	<0.0001
cough	18	4.48	0.0094
cold	82	1.45	0.4154
medicin	9	11.20	<0.0001
fever	13	26.20	<0.0001

Table 1: Keyword effects.

tion regarding the user’s health. We also sample a randomly selected set of 1609 tweets from times when the users did not have influenza as a control. We find 58 tweets from 17 (17/35 = 48.57%) individuals in our study that are about the user being sick. We also find zero tweets about ILI during times when they did *not* have influenza. While the use of a “human” classifier clearly does not scale, it allows for an approximately 100.0% accurate classification. Since regular machine learning algorithms perform much worse than 100.0% accuracy, the human classifier gives us an upper limit to the accuracy of a health monitoring system based off of tweet classification (see table 2.)

4. FREQUENCY BASED SIGNALS

It is likely that a user tweets at a different rate when she is ill than she normally does. To detect this, we perform one-dimensional anomaly detection on each user’s monthly tweeting rate as follows. First, we calculate the number of tweets in each month in the study period and discard any months where the user tweets less than ten times. This avoids issues caused by the user starting or stopping their use of Twitter. We then calculate the z-score of the tweeting rate of the month that the user is ill by

Sick	Not Sick	
17	18	Sick
0	66	Not Sick

Table 2: Confusion matrix of a Tweet-Classification based diagnosis system. Rows are of true values, columns are of predicted values.

Sick	Not Sick	
14	25	Sick
27	192	Not Sick

Table 3: Confusion matrix of the classifier based on anomalous tweeting rates. Rows are of true values, columns are of predicted values.

$$z = \frac{|x - \bar{x}|}{\hat{s}} \quad (1)$$

Where \bar{x} and \hat{s} are the estimated mean and standard deviation of the user’s tweeting rate. [6] We repeat this process for months when the user is not sick. We then classify the user as sick if $z > 1.411$ where 1.411 was chosen through leave one out cross validation. We find a significant difference between the z-scores for months when a user had influenza and months when the user did not ($p = 0.01303$, two-sample Kolmogorov-Smirnov test). Most of the time individuals are not sick (219 / 258 = 84.88%), resulting in a highly biased sample. Thus we optimize based on the F_1 score instead of accuracy. The optimal z-score cutoff results in $F_1 = 35.0\%$ and an area under the ROC curve of . (See table 3.)

5. NETWORK BASED SIGNALS

Even if a user is not currently active at all on Twitter, her friends or followers may give clues to her health status. We consider all text that a user’s friends or followers tweeted and perform similar keyword analysis as was applied to the user’s tweets in section 3. Because the number and activity of friends or followers greatly varies regardless of a user’s health status, we normalize the counts here by the total number of characters her friends or followers tweeted. We find that most of the tested classifiers are able to detect a signal in both the user’s followers’ and friends’ streams (see figure 3.)

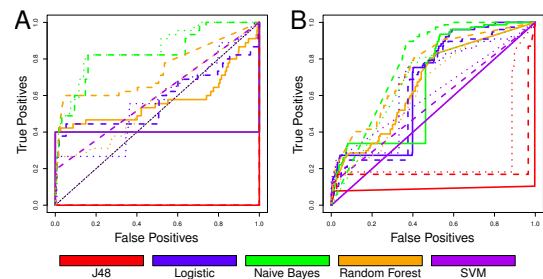


Figure 3: The ROC of classifiers based off Tweets from (a) accounts that follow a user and (b) accounts that a user follows. Line coloring and style are equivalent to figure 2.

We further analyse the strength of these classifiers by generating 100 of each classifier using 10 fold cross validation and calculate their area under the ROC curve. We then perform an analysis of variance test based off of the source of the data (friends or followers), the number of keywords used and the classifier’s algorithm (see table 4.) We find that the choice in classifier and the length of the feature vector have a significant effect on strength. More interestingly, we

	Df	Sum Sq	F value	Pr(>F)
Source	1	107.16	1290.82	<2e-16
Keyword Size	1	72.19	869.66	<2e-16
Classifier	3	752.55	3021.61	<2e-16
Residuals	109194	9602		

Table 4: Results from an analysis of variance of the area under the ROC curve for classifiers based on tweets from an individual’s social network. Factors are whether the data is from the user’s friends or followers, the number of keywords chosen and the classifier.

Classifier	Area under ROC	Accuracy
AdaBoost	.9961	99.53
Bayesian	.9078	92.08
Decision Tree	.9877	99.22
Logit Boost	.9986	99.22
Weighted Voting	.9783	93.17
Baseline	.8544	89.72

Table 5: Performance of the meta classifiers. The presented baseline is the classifier based on datamined keywords – the highest performing individual classifier.

find that classifiers that use tweets from accounts that follow the user are significantly better at diagnosing the user than classifiers that use tweets from accounts that the user follows.

6. META CLASSIFIER

So far we have considered five separate methods for detecting illness based off of a user’s Twitter activity. However, there is no reason that we cannot combine these methods to get a stronger signal. For example, while mining the user’s text is the best of the five methods, she may stop tweeting while sick, which would be detected by the frequency-based anomaly classifier. Aggregating multiple classifiers by a ‘meta-classifier’ has been shown to be an effective method for increasing classification accuracy. [15, 14]

We start by selecting the classifier from each of the previous five approaches that has the largest area under the ROC curve (see figure 4.a.) We then use the predicted distributions from these classifiers as the feature vector for the meta classifier. We use Ada Boost, Bayesian classification, J48 decision trees, logit boost, and weighted voting to evaluate the meta-dataset. We then evaluate these methods with leave-one-out cross validation and see an increase in area under ROC and accuracy compared to the best individual classifier (see figure 4.b.)

7. CONCLUSIONS

In this paper, we have shown that it is possible to diagnose an individual based off of her social media data with high accuracy. Computational approaches to aid in disease diagnosis has been approached before, however they have been developed with a medical setting in mind. That is, the question addressed was “can we diagnose an individual based off data gathered from medical tests run on her?” instead of “can we diagnose an individual solely based off of




Figure 4: The accuracy of the previous classifiers (a) and the accuracy of various classifiers that use the previous classifier’s results as features (b).

publicly available social media data?” While we focus on the relatively benign case of remotely reconstructing a confidential diagnosis of influenza, these methods could also be applied to stigmatized diseases, such as HIV, where being able to determine if an individual is HIV positive without her knowledge and with only her Twitter handle could result in serious social or economic effects. Half of the users explicitly stated that they were sick, and we were able to confidently determine illness in the other half of the cases through their data. It would seem that simply avoiding discussing a illness is not enough to hide one’s health in the age of big data.

8. ACKNOWLEDGMENTS

Marcel Salathé received funding through a Branco Weiss fellowship. We thank the Pennsylvania State University’s Student Health Center for aid in collecting data. We thank Brian Lambert and Cosme Adrover Pacheco for their valuable comments on the paper. The machine learning tool set Weka version 3.7.10 was used in this paper [7].

9. REFERENCES

- [1] T. Bodnar and M. Salathé. Validating models for disease detection using twitter. In *Proceedings of the 22Nd International Conference on World Wide Web Companion*, WWW ’13 Companion, pages 699–702, Republic and Canton of Geneva, Switzerland, 2013. International World Wide Web Conferences Steering Committee.
- [2] D. Butler. When Google got flu wrong. *Nature*, 494(7436):155–156, Feb. 2013.
- [3] E. H. Chan, T. F. Brewer, L. C. Madoff, M. P. Pollack, A. L. Sonrick, M. Keller, C. C. Freifeld, M. Blench, A. Mawudeku, and J. S. Brownstein. Global capacity for emerging infectious disease detection. *Proceedings of the National Academy of Sciences*, 107(50):21701–21706, 2010.
- [4] A. Culotta. Towards detecting influenza epidemics by analyzing Twitter messages. In *the First Workshop*, pages 115–122, New York, New York, USA, 2010. ACM Press.
- [5] S. Goel, J. M. Hofman, S. Lahaie, D. M. Pennock, and D. J. Watts. Predicting consumer behavior with Web search. *Proceedings of the National Academy of Sciences of the United States of America*, 107(41):17486–17490, Oct. 2010.
- [6] F. E. Grubbs. Procedures for Detecting Outlying

- Observations in Samples. *Technometrics*, 11(1):1–21, 1969.
- [7] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten. The weka data mining software: An update. *SIGKDD Explor. Newsl.*, 11(1):10–18, Nov. 2009.
- [8] D. L. Heymann and G. R. Rodier. Hot spots in a wired world: {WHO} surveillance of emerging and re-emerging infectious diseases. *The Lancet Infectious Diseases*, 1(5):345 – 353, 2001.
- [9] A. Lamb, M. J. Paul, and M. Dredze. Separating fact from fear: Tracking flu infections on twitter. In *Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 789–795, Atlanta, Georgia, June 2013. Association for Computational Linguistics.
- [10] R. L. Marquet, A. I. Bartelds, S. P. van Noort, C. E. Koppeschaar, J. Paget, F. G. Schellevis, and J. van der Zee. Internet-based monitoring of influenza-like illness (ILI) in the general population of the Netherlands during the 2003-2004 influenza season. *BMC public health*, 6(1):242, 2006.
- [11] D. R. Olson, K. J. Konty, M. Paladini, C. Viboud, and L. Simonsen. Reassessing Google Flu Trends Data for Detection of Seasonal and Pandemic Influenza: A Comparative Epidemiological Study at Three Geographic Scales. *PLoS computational biology*, 9(10):e1003256, Oct. 2013.
- [12] M. F. Porter. An algorithm for suffix stripping. *Program: electronic library and information systems*, 14(3):130–137, 1980.
- [13] M. Salathé, L. Bengtsson, T. J. Bodnar, D. D. Brewer, J. S. Brownstein, C. Buckee, E. M. Campbell, C. Cattuto, S. Khandelwal, P. L. Mabry, and A. Vespignani. Digital epidemiology. *PLoS computational biology*, 8(7):e1002616, Jul 2012.
- [14] L. Todorovski and S. Džeroski. Combining classifiers with meta decision trees. *Mach. Learn.*, 50(3):223–249, Mar. 2003.
- [15] G. Tsirogiannis, D. Frossyniotis, K. Nikita, and A. Stafylopatis. A meta-classifier approach for medical diagnosis. In G. Vouros and T. Panayiotopoulos, editors, *Methods and Applications of Artificial Intelligence*, volume 3025 of *Lecture Notes in Computer Science*, pages 154–163. Springer Berlin Heidelberg, 2004.
- [16] S. P. van Noort, M. Muehlen, H. Rebelo de Andrade, C. Koppeschaar, J. M. Lima Lourenco, and M. G. Gomes. Gripenet: an internet-based system to monitor influenza-like illness uniformly across Europe. *Euro Surveill.*, 12(7):5–6, Jul 2007.
- [17] P. Willett. The Porter stemming algorithm: then and now. *Program: electronic library and information systems*, 2006.

APPENDIX

A. KEYWORD RECOMMENDATIONS

While our system should be trusted more than one based simply off of aggregated tweets, it is more computationally intensive than simply pulling data from a keyword stream. These systems require the user to select a specific set of keywords before data collection can begin. Keywords representing symptoms such as “flu”, “cough”, “sore throat”, and “headache” are often chosen. We suggest that the thirty² keywords with the highest positive predictive value (see table 6) be chosen as the parameters for a keyword stream. In addition to keywords related to symptoms (e.g. “flu” or “sick”) we also find keywords related to treatments (e.g. “health,” “prayer” or “recovery”) and keywords related to negative mood (e.g. vulgarities) to be more commonly tweeted when a user is ill.

Keyword	Ratio	cont.	
flu	34.424	walk	3.077
health	11.360	childr	6.820
sick	5.019	incred	6.820
track	10.952	meal	6.820
stud	3.508	longer	6.820
asshol	9.090	succes	26.765
ton	9.090	accis	26.765
particip	20.667	holida	26.765
salt	20.667	luv	26.765
recov	40.118	oblig	26.765
fuck	2.963	path	26.764
sham	13.64	pract	26.764
row	10.180	prayer	26.765
win	2.947	reserv	26.765
rt	3.077	riot	26.765

Table 6: The thirty keyword stems with the highest positive predictive power ranked by significance. Ratio is calculated as the rate of occurrence when a user is sick over the rate when a user is not sick.

²The Twitter API limits queries to thirty keywords.